
Recommender Systems with Python

Welcome to the code notebook for Recommender Systems with Python. In this lecture we will develop basic recommendation systems using Python and pandas.

In this notebook, we will focus on providing a basic recommendation system by suggesting items that are most similar to a particular item, in this case, movies. Keep in mind, this is not a true robust recommendation system, to describe it more accurately, it just tells you what movies/items are most similar to your movie choice.

There is no project for this topic, instead you have the option to work through the advanced lecture version of this notebook (totally optional!).

Let's get started!

Import Libraries

```
In [1]: import numpy as np
import pandas as pd
```

Get the Data

```
In [2]: column_names = ['user_id', 'item_id', 'rating', 'timestamp']
df = pd.read_csv("u.data", sep='\t', names=column_names)
```

```
In [3]: df.head()
```

Out[3]:

| | user_id | item_id | rating | timestamp |
|---|---------|---------|--------|-----------|
| 0 | 0 | 50 | 5 | 881250949 |
| 1 | 0 | 172 | 5 | 881250949 |
| 2 | 0 | 133 | 1 | 881250949 |
| 3 | 196 | 242 | 3 | 881250949 |
| 4 | 186 | 302 | 3 | 891717742 |

Now let's get the movie titles:

```
In [4]: movie_titles = pd.read_csv("Movie_Id_Titles (1).txt")  
movie_titles.head()
```

Out[4]:

| | item_id | title |
|---|---------|-------------------|
| 0 | 1 | Toy Story (1995) |
| 1 | 2 | GoldenEye (1995) |
| 2 | 3 | Four Rooms (1995) |
| 3 | 4 | Get Shorty (1995) |
| 4 | 5 | Copycat (1995) |

We can merge them together:

```
In [5]: df = pd.merge(df, movie_titles, on='item_id')
df.head()
```

```
Out[5]:
```

| | user_id | item_id | rating | timestamp | title |
|---|---------|---------|--------|-----------|------------------|
| 0 | 0 | 50 | 5 | 881250949 | Star Wars (1977) |
| 1 | 290 | 50 | 5 | 880473582 | Star Wars (1977) |
| 2 | 79 | 50 | 4 | 891271545 | Star Wars (1977) |
| 3 | 2 | 50 | 5 | 888552084 | Star Wars (1977) |
| 4 | 8 | 50 | 5 | 879362124 | Star Wars (1977) |

EDA

Let's explore the data a bit and get a look at some of the best rated movies.

Visualization Imports

```
In [6]: import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('white')
%matplotlib inline
```

Let's create a ratings dataframe with average rating and number of ratings:

```
In [7]: df.groupby('title')['rating'].mean().sort_values(ascending=False).head()
```

```
Out[7]: title
They Made Me a Criminal (1939)      5.0
Marlene Dietrich: Shadow and Light (1996)  5.0
Saint of Fort Washington, The (1993)      5.0
Someone Else's America (1995)          5.0
Star Kid (1997)                      5.0
Name: rating, dtype: float64
```

```
In [8]: df.groupby('title')['rating'].count().sort_values(ascending=False).head()
```

```
Out[8]: title
Star Wars (1977)          584
Contact (1997)            509
 Fargo (1996)             508
Return of the Jedi (1983)  507
Liar Liar (1997)          485
Name: rating, dtype: int64
```

```
In [9]: ratings = pd.DataFrame(df.groupby('title')['rating'].mean())
ratings.head()
```

```
Out[9]:
```

| | rating |
|---------------------------|----------|
| title | |
| 'Til There Was You (1997) | 2.333333 |
| 1-900 (1994) | 2.600000 |
| 101 Dalmatians (1996) | 2.908257 |
| 12 Angry Men (1957) | 4.344000 |
| 187 (1997) | 3.024390 |

Now set the number of ratings column:

```
In [10]: ratings['num of ratings'] = pd.DataFrame(df.groupby('title')['rating'].count())
ratings.head()
```


Out[10]:

| | rating | num of ratings |
|---------------------------|----------|----------------|
| title | | |
| 'Til There Was You (1997) | 2.333333 | 9 |
| 1-900 (1994) | 2.600000 | 5 |
| 101 Dalmatians (1996) | 2.908257 | 109 |
| 12 Angry Men (1957) | 4.344000 | 125 |
| 187 (1997) | 3.024390 | 41 |

Now a few histograms:


```
In [11]: plt.figure(figsize=(10,4))  
ratings['num of ratings'].hist(bins=70)
```

Out[11]: <Axes: >




```
In [12]: plt.figure(figsize=(10,4))  
ratings['rating'].hist(bins=70)
```

Out[12]: <Axes: >



```
In [13]: sns.jointplot(x='rating',y='num of ratings',data=ratings,alpha=0.5)
```

```
Out[13]: <seaborn.axisgrid.JointGrid at 0x1ba5d701210>
```



Okay! Now that we have a general idea of what the data looks like, let's move on to creating a simple recommendation system:

Recommending Similar Movies

Now let's create a matrix that has the user ids on one axis and the movie title on another axis. Each cell will then consist of the rating the user gave to that movie. Note there will be a lot of NaN values, because most people have not seen most of the movies.

```
In [14]: moviemat = df.pivot_table(index='user_id',columns='title',values='rating')
moviemat.head()
```

Out[14]:

| | 'Til There Was You (1997) | 1-900 (1994) | 101 Dalmatians (1996) | 12 Angry Men (1957) | 187 (1997) | 2 Days in the Valley (1996) | 20,000 Leagues Under the Sea (1954) | 2001: A Space Odyssey (1968) | 3 Ninjas: High Noon At Mega Mountain (1998) | 39 Steps, The (1935) | ... | Yankee Zulu (1994) | Year of the Horse (1997) | You So Crazy (1994) | Young Frankenstein (1974) |
|---------|---------------------------------------|-----------------|-----------------------------|------------------------------|---------------|---|---|---------------------------------------|--|-------------------------------|-----|--------------------------|-----------------------------------|------------------------------|---------------------------------|
| user_id | | | | | | | | | | | | | | | |
| 0 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | NaN |
| 1 | NaN | NaN | 2.0 | 5.0 | NaN | NaN | 3.0 | 4.0 | NaN | NaN | ... | NaN | NaN | NaN | 5.0 |
| 2 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | 1.0 | NaN | ... | NaN | NaN | NaN | NaN |
| 3 | NaN | NaN | NaN | NaN | 2.0 | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | NaN |
| 4 | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | NaN | ... | NaN | NaN | NaN | NaN |

5 rows × 1664 columns

Most rated movie:

```
In [15]: ratings.sort_values('num of ratings',ascending=False).head(10)
```

Out[15]:

| | rating | num of ratings |
|-------------------------------|----------|----------------|
| title | | |
| Star Wars (1977) | 4.359589 | 584 |
| Contact (1997) | 3.803536 | 509 |
| Fargo (1996) | 4.155512 | 508 |
| Return of the Jedi (1983) | 4.007890 | 507 |
| Liar Liar (1997) | 3.156701 | 485 |
| English Patient, The (1996) | 3.656965 | 481 |
| Scream (1996) | 3.441423 | 478 |
| Toy Story (1995) | 3.878319 | 452 |
| Air Force One (1997) | 3.631090 | 431 |
| Independence Day (ID4) (1996) | 3.438228 | 429 |

Let's choose two movies: starwars, a sci-fi movie. And Liar Liar, a comedy.

```
In [16]: ratings.head()
```

Out[16]:

| | rating | num of ratings |
|---------------------------|----------|----------------|
| title | | |
| 'Til There Was You (1997) | 2.333333 | 9 |
| 1-900 (1994) | 2.600000 | 5 |
| 101 Dalmatians (1996) | 2.908257 | 109 |
| 12 Angry Men (1957) | 4.344000 | 125 |
| 187 (1997) | 3.024390 | 41 |

Now let's grab the user ratings for those two movies:

```
In [17]: starwars_user_ratings = moviemat['Star Wars (1977)']
liarliar_user_ratings = moviemat['Liar Liar (1997)']
starwars_user_ratings.head()
```

```
Out[17]: user_id
0      5.0
1      5.0
2      5.0
3      NaN
4      5.0
Name: Star Wars (1977), dtype: float64
```

We can then use `corrwith()` method to get correlations between two pandas series:

```
In [18]: similar_to_starwars = moviemat.corrwith(starwars_user_ratings)
similar_to_liarliar = moviemat.corrwith(liarliar_user_ratings)
```

```
D:\anaconda\Lib\site-packages\numpy\lib\function_base.py:2846: RuntimeWarning: Degrees of freedom <= 0 for slice
```

```
    c = cov(x, y, rowvar, dtype=dtype)
```

```
D:\anaconda\Lib\site-packages\numpy\lib\function_base.py:2705: RuntimeWarning: divide by zero encountered in divide
```

```
    c *= np.true_divide(1, fact)
```

Let's clean this by removing NaN values and using a DataFrame instead of a series:

```
In [19]: corr_starwars = pd.DataFrame(similar_to_starwars, columns=['Correlation'])
corr_starwars.dropna(inplace=True)
corr_starwars.head()
```

Out[19]:

| | Correlation |
|---------------------------|-------------|
| title | |
| 'Til There Was You (1997) | 0.872872 |
| 1-900 (1994) | -0.645497 |
| 101 Dalmatians (1996) | 0.211132 |
| 12 Angry Men (1957) | 0.184289 |
| 187 (1997) | 0.027398 |

Now if we sort the dataframe by correlation, we should get the most similar movies, however note that we get some results that don't really make sense. This is because there are a lot of movies only watched once by users who also watched star wars (it was the most popular movie).

```
In [20]: corr_starwars.sort_values('Correlation',ascending=False).head(10)
```

Out[20]:

| | Correlation |
|---|-------------|
| title | |
| Commandments (1997) | 1.0 |
| Cosi (1996) | 1.0 |
| No Escape (1994) | 1.0 |
| Stripes (1981) | 1.0 |
| Man of the Year (1995) | 1.0 |
| Hollow Reed (1996) | 1.0 |
| Beans of Egypt, Maine, The (1994) | 1.0 |
| Good Man in Africa, A (1994) | 1.0 |
| Old Lady Who Walked in the Sea, The (Vieille qui marchait dans la mer, La) (1991) | 1.0 |
| Outlaw, The (1943) | 1.0 |

Let's fix this by filtering out movies that have less than 100 reviews (this value was chosen based off the histogram from earlier).

```
In [21]: corr_starwars = corr_starwars.join(ratings['num of ratings'])
corr_starwars.head()
```

Out[21]:

| | Correlation | num of ratings |
|---------------------------|-------------|----------------|
| title | | |
| 'Til There Was You (1997) | 0.872872 | 9 |
| 1-900 (1994) | -0.645497 | 5 |
| 101 Dalmatians (1996) | 0.211132 | 109 |
| 12 Angry Men (1957) | 0.184289 | 125 |
| 187 (1997) | 0.027398 | 41 |

Now sort the values and notice how the titles make a lot more sense:

```
In [22]: corr_starwars[corr_starwars['num of ratings']>100].sort_values('Correlation',ascending=False).head()
```

Out[22]:

| | Correlation | num of ratings |
|--|-------------|----------------|
| title | | |
| Star Wars (1977) | 1.000000 | 584 |
| Empire Strikes Back, The (1980) | 0.748353 | 368 |
| Return of the Jedi (1983) | 0.672556 | 507 |
| Raiders of the Lost Ark (1981) | 0.536117 | 420 |
| Austin Powers: International Man of Mystery (1997) | 0.377433 | 130 |

Now the same for the comedy Liar Liar:

```
In [23]: corr_liarliar = pd.DataFrame(similar_to_liarliar,columns=['Correlation'])
corr_liarliar.dropna(inplace=True)
corr_liarliar = corr_liarliar.join(ratings['num of ratings'])
corr_liarliar[corr_liarliar['num of ratings']>100].sort_values('Correlation',ascending=False).head()
```

Out[23]:

| | Correlation | num of ratings |
|-----------------------|-------------|----------------|
| title | | |
| Liar Liar (1997) | 1.000000 | 485 |
| Batman Forever (1995) | 0.516968 | 114 |
| Mask, The (1994) | 0.484650 | 129 |
| Down Periscope (1996) | 0.472681 | 101 |
| Con Air (1997) | 0.469828 | 137 |

Type *Markdown* and LaTeX: α^2

In []:

In []: